

Diagnosing and Fixing Prompt Unfaithfulness in SAM-Med3D via Paired Prompt-Contrastive Training

Junqing Liu

Abstract

3D medical Segment Anything Models promise interactive segmentation via point prompts, yet we discover they suffer from *prompt unfaithfulness*: the model largely ignores prompt location and predicts based on image priors alone, achieving only 20.8% Switch Accuracy—the rate at which predictions correctly change when prompts target different organs. This failure is distinct from prompt robustness (only 1.3-point Dice drop under 10-voxel noise) and is hidden by standard Dice evaluation. We identify the root cause as single-prompt training, which never requires the model to produce different outputs for different prompts on the same image. To address this, we propose *Paired Prompt-Contrastive Training*: for each image, we sample two organ prompts and supervise both predictions through a shared LoRA-adapted decoder. Comprehensive mechanism ablation reveals that paired prompt exposure itself is the dominant factor (+40 points Switch Accuracy), while auxiliary losses (grounding regularization, switch margin) provide only marginal additional benefit (+0–5 points). Our method achieves Switch Accuracy 0.993 ± 0.012 (from 0.208 baseline) with only 4.4-point Dice cost on AMOS across three seeds, generalizes to BTCV external validation (Switch 1.000), and adds merely 5.3% training overhead with zero inference cost. Code will be released upon acceptance.

1 Introduction

Segment Anything Models (SAM) adapted for 3D medical imaging [Wang et al., 2023, Du et al., 2023, Ma et al., 2024, Zhu et al., 2024] enable interactive organ segmentation through point prompts: a clinician clicks on a target organ, and the model produces a volumetric segmentation mask. This paradigm assumes a fundamental contract—changing the prompt target should change the predicted mask. The promise of such models lies precisely in this interactive controllability, which enables efficient annotation workflows and intraoperative guidance.

We discover that this contract is systematically violated. When evaluated on multi-organ abdominal CT, SAM-Med3D [Wang et al., 2023] achieves only 20.8% *Switch Accuracy*—the rate at which the model correctly changes its prediction when the prompt moves from one organ to another (Figure 1). The model predicts the same dominant organ regardless of prompt location, effectively ignoring user input. This means that in nearly 80% of cases, clicking on a different organ produces the same segmentation output—a catastrophic failure for interactive clinical use.

This failure is *not* a robustness problem. Recent work [Chattopadhyay et al., 2026] studies sensitivity to prompt perturbations, finding that models degrade under

noise. In contrast, SAM-Med3D is remarkably robust to prompt noise (only 1.3-point Dice drop at 10-voxel perturbation). The failure we identify is a *controllability* problem: the model has learned strong anatomical priors that override prompt information entirely. Standard Dice evaluation completely hides this problem because Dice rewards accurate segmentation regardless of whether the model actually follows the prompt.

We trace the root cause to single-prompt training. When each training step presents only one prompt per image, the model can achieve low segmentation loss by exploiting image priors—always predicting the largest or most salient organ—without learning prompt-conditioned behavior. The model never needs to demonstrate that different prompts produce different outputs. This is analogous to a language model that produces the same response regardless of the input query: technically functional but fundamentally uncontrollable.

Our fix is conceptually simple: *Paired Prompt-Contrastive Training*. For each training image, we sample two organ prompts, compute predictions for both using a shared decoder, and supervise both outputs. This forces the model to produce different masks for different prompts on the same image—it can no longer satisfy the training objective by ignoring prompts. Through comprehensive mechanism ablation, we discover that this paired exposure is the dominant factor—achieving +40 points Switch Accuracy improvement—while auxiliary losses (grounding regularization, switch margin) provide only marginal additional benefit (+0–5 points).

Our contributions are:

1. We identify *prompt unfaithfulness* as a distinct failure mode in 3D medical SAM models, measurable via Switch Accuracy, and show it is hidden by standard Dice evaluation yet catastrophic for interactive use.
2. We propose Paired Prompt-Contrastive Training and demonstrate through factor-isolation ablation that paired prompt exposure—not any specific loss function—is the core mechanism restoring faithfulness.
3. We validate across datasets (AMOS, BTCV) with 3-seed experiments, achieving Switch Accuracy 0.993 ± 0.012 with Dice 0.676 ± 0.006 , at only 5.3% training overhead and zero inference cost.
4. We show that prompt unfaithfulness is specific to 3D medical SAM adaptation (SAM2 achieves 0.954 Switch on natural images), suggesting it arises from limited medical training data and strong anatomical priors.




Figure 1: Prompt unfaithfulness in 3D medical SAM. (a) Expected behavior: different point prompts should produce different segmentation masks. (b) Observed behavior: SAM-Med3D produces nearly identical predictions regardless of prompt target (Switch Accuracy = 20.8%). (c) After our Paired Prompt-Contrastive Training: the model correctly switches predictions to match the prompted organ (Switch Accuracy = 99.3%).

2 Related Work

Foundation Models for 3D Medical Segmentation. The Segment Anything Model (SAM) [Kirillov et al., 2023] introduced promptable segmentation for natural images, demonstrating remarkable zero-shot generalization through point, box, and mask prompts. SAM2 [Ravi et al., 2024] extended this paradigm to video with memory-augmented tracking, enabling temporal consistency across frames. The adaptation of these models to 3D medical imaging has become a major research direction.

SAM-Med3D [Wang et al., 2023] trains a volumetric SAM architecture on large-scale CT data with point prompts, achieving competitive performance on multi-organ segmentation. MedSAM [Ma et al., 2024] fine-tunes the original SAM for medical bounding-box prompts across diverse imaging modalities. Medical SAM2 [Zhu et al., 2024] re-frames 3D medical segmentation as video object tracking via SAM2, introducing a self-sorting memory bank for cross-slice consistency. MA-SAM [Chen et al., 2024a] proposes modality-agnostic adapters that enable cross-modal transfer without modality-specific training. SAM2-Adapter [Chen et al., 2024b] introduces task-specific adapters to overcome SAM2 limitations in medical, camouflage, and shadow segmentation domains.

A common thread across these works is evaluation via segmentation accuracy metrics (Dice, IoU) without assessing whether the model faithfully follows prompt targets. Our work reveals that high Dice scores can coexist with catastrophic prompt unfaithfulness—a gap that standard evaluation completely misses.

Parameter-Efficient Adaptation for Medical Imaging. LoRA [Hu et al., 2022] enables efficient fine-tuning by injecting trainable low-rank matrices into frozen model layers, dramatically reducing the number of trainable parameters while maintaining adaptation quality. This approach has been

widely adopted in medical imaging.

LoRA-MedSAM [Wu et al., 2024] applies LoRA to MedSAM, reducing fine-tuning time by 22.7% while improving Dice scores. FATE-SAM [Deng et al., 2025] achieves few-shot 3D segmentation by reassembling pre-trained SAM2 modules without any gradient-based training, leveraging support examples for prompt-free inference. tCURLoRA [Cheng et al., 2025] proposes tensor CUR decomposition as an alternative to standard LoRA, achieving better parameter efficiency for medical segmentation tasks.

Our work uses standard LoRA adaptation but reveals a crucial insight: the *training paradigm* (paired vs. single prompt) matters far more than the adaptation architecture or parameter count. Even with identical LoRA configurations, single-prompt training produces unfaithful models while paired-prompt training produces faithful ones.

Prompt Robustness in Medical Segmentation. Recent work has begun examining how medical segmentation models respond to imperfect prompts. Chattopadhyay et al. [2026] systematically evaluate how foundational 3D medical segmentation models respond to controlled perturbations of dense visual prompts, revealing that models rely heavily on spatial shape cues and exhibit varying resilience to different perturbation types. Huang et al. [2024] study perturbed prompts for robust SAM adaptation in polyp segmentation, proposing augmentation strategies that improve stability. Li et al. [2024] develop saliency-guided prompt distillation to maintain segmentation quality under noisy or generic prompts in clinical workflows.

These works study *robustness*—stability of predictions under small input perturbations. We identify a fundamentally different and more severe failure: *faithfulness*—whether the model responds to target changes at all. The distinction is critical: a robust-but-unfaithful model produces stable predictions that consistently ignore the user, while a faithful

model correctly changes its output when the user indicates a different target. SAM-Med3D is robust (1.3pt Dice drop at 10-voxel shift) but unfaithful (Switch Accuracy 20.8%)—demonstrating that these are orthogonal properties requiring fundamentally different solutions.

Contrastive and Paired Training in Medical Analysis.

Contrastive learning has been extensively applied to medical image analysis for representation learning and pre-training [Zhao et al., 2024]. Recent advances address specific challenges: Liu et al. [2024] identify and overcome dimensional collapse in self-supervised contrastive learning for medical images, showing that high inter-image similarity in medical data causes representation degeneration. Wang et al. [2025] propose domain and task-focused example selection for data-efficient contrastive medical segmentation.

Our paired prompt training shares the principle of exposing the model to contrasting conditions within a single optimization step. However, it differs fundamentally in objective: rather than learning invariant or discriminative *representations*, we enforce prompt-conditioned output *differentiation*—the model must produce different segmentation masks for different prompts. This is a supervised, task-level constraint rather than a representation-level one, and it directly addresses the prompt faithfulness failure rather than general feature quality.

3 Method

3.1 Problem Formulation

Let f_θ be a promptable segmentation model that takes a 3D image $\mathbf{x} \in \mathbb{R}^{D \times H \times W}$ and a point prompt $p = (d, h, w)$ to produce a binary mask $f_\theta(\mathbf{x}, p) \in \{0, 1\}^{D \times H \times W}$. We define *prompt faithfulness* as the property that changing the prompt target changes the prediction accordingly.

Switch Accuracy. Given an image $\mathbf{x}$ containing organs A and B with ground-truth masks y_A and y_B , and prompts $p_A \in A, p_B \in B$, we define a *switch* as successful if both conditions hold simultaneously:

$$\text{Dice}(f_\theta(\mathbf{x}, p_A), y_A) > \text{Dice}(f_\theta(\mathbf{x}, p_A), y_B) \quad (1)$$

$$\text{Dice}(f_\theta(\mathbf{x}, p_B), y_B) > \text{Dice}(f_\theta(\mathbf{x}, p_B), y_A) \quad (2)$$

Switch Accuracy is the fraction of paired prompts satisfying both conditions across a fixed evaluation bank. This metric captures whether the model produces the *correct* organ for each prompt, not merely whether it produces *something*.

3.2 Why Standard Training Fails

Standard training optimizes $\mathcal{L}_{\text{seg}}(f_\theta(\mathbf{x}, p), y)$ independently for each prompt. This objective is satisfiable without prompt faithfulness: a model that always predicts the largest organ achieves reasonable Dice regardless of prompt location. Formally, if organ A dominates the volume, then $f_\theta(\mathbf{x}, p) \approx y_A$ for all p achieves high average Dice without ever using p . The

loss contains no term enforcing that $f_\theta(\mathbf{x}, p_A) \neq f_\theta(\mathbf{x}, p_B)$ when $A \neq B$.

This explains why standard LoRA fine-tuning (seg-only) improves Dice from 0.450 to 0.720 but Switch Accuracy only reaches 0.483—the model learns better segmentation but not better prompt following.

3.3 Paired Prompt-Contrastive Training

We address this by training with paired prompts (Figure 2). For each image $\mathbf{x}$ containing organs A and B :

1. Compute image embedding $\mathbf{e} = \text{Enc}(\mathbf{x})$ once (frozen encoder)
2. Decode with prompt p_A : $\hat{y}_A = \text{Dec}_\phi(\mathbf{e}, p_A)$
3. Decode with prompt p_B : $\hat{y}_B = \text{Dec}_\phi(\mathbf{e}, p_B)$
4. Supervise both predictions against respective ground truths

The decoder Dec_ϕ (prompt encoder + mask decoder) is shared between both passes and adapted via LoRA [Hu et al., 2022] with rank 8, yielding 1.9M trainable parameters (1.89% of total 100.5M). The image encoder remains frozen, so the second forward pass reuses the cached embedding $\mathbf{e}$ with minimal overhead.

Loss Function. The total loss combines three terms:

$$\mathcal{L} = \mathcal{L}_{\text{seg}} + \lambda_s \mathcal{L}_{\text{switch}} + \lambda_g \mathcal{L}_{\text{ground}} \quad (3)$$

The segmentation loss averages Dice+BCE on both predictions:

$$\mathcal{L}_{\text{seg}} = \frac{1}{2} [\ell(\hat{y}_A, y_A) + \ell(\hat{y}_B, y_B)] \quad (4)$$

The switch margin loss enforces output separation:

$$\begin{aligned} \mathcal{L}_{\text{switch}} = & [m + D(\hat{y}_A, y_B) - D(\hat{y}_A, y_A)]_+ \\ & + [m + D(\hat{y}_B, y_A) - D(\hat{y}_B, y_B)]_+ \end{aligned} \quad (5)$$

where $D(\cdot, \cdot)$ is soft Dice similarity, $m=0.2$ is the margin, and $[\cdot]_+ = \max(0, \cdot)$. This loss penalizes predictions that are more similar to the wrong organ than to the correct one, maintaining a persistent gradient signal throughout training.

The grounding loss ensures prompt points lie inside predictions:

$$\mathcal{L}_{\text{ground}} = -\log \sigma(\hat{y}_A[p_A]) - \log \sigma(\hat{y}_B[p_B]) \quad (6)$$

where σ is the sigmoid function and $\hat{y}[p]$ denotes the logit at the prompt location.

We use $\lambda_s=1.0$ and $\lambda_g=0.5$ throughout. As our ablation demonstrates (§4.4), the specific loss weights matter far less than the paired training paradigm itself.




Figure 2: Paired Prompt-Contrastive Training framework. The frozen image encoder computes a shared embedding. Two different organ prompts are decoded through the same LoRA-adapted decoder, producing two predictions supervised by their respective ground truths. The switch margin loss enforces that each prediction is closer to its target organ than to the alternative.

Key Insight. The critical mechanism is not any particular loss function but the *paired prompt exposure*: by presenting two different prompts for the same image in a single training step, the model cannot satisfy the objective by ignoring prompts. It must learn to condition its output on the prompt to minimize loss on both predictions simultaneously.

4 Experiments

4.1 Experimental Setup

Datasets. We use AMOS [Ji et al., 2022] (240 training, 60 validation CT volumes, 13 abdominal organs, resampled to 1.5mm isotropic, 128^3 crops) as the primary dataset and BTCV [Landman et al., 2015] (21 training, 9 validation, 13 organs) for external validation. Both datasets provide multi-organ annotations enabling paired prompt sampling.

Model and Adaptation. SAM-Med3D-turbo [Wang et al., 2023] with frozen ViT-B image encoder. LoRA adaptation (rank=8) on prompt encoder and mask decoder, yielding 1.9M trainable parameters (1.89% of 100.5M total).

Training. 50 epochs, AdamW optimizer (LR 10^{-4} , weight decay 0.01), 2-epoch linear warmup followed by cosine decay to 10^{-6} , batch size 1 with gradient accumulation 4. Three random seeds (42, 123, 7) for main experiments. Checkpoint selection: maximum validation Dice subject to Switch Accuracy ≥ 0.90 .

Evaluation Protocol. We construct a fixed prompt bank of 514 paired prompts (deterministic across all methods) to ensure fair comparison. Each pair consists of two point prompts targeting different organs in the same volume. We report Dice

Table 1: Main results on AMOS and BTCV. Our method achieves near-perfect Switch Accuracy with modest Dice cost. McNemar test: $p < 0.001$ for Switch improvement across all seeds.

Method	Dice	Switch Acc.
<i>AMOS (3 seeds, mean $\pm$ std)</i>		
Frozen SAM-Med3D	0.450	0.208
seg-only (LoRA)	0.720 ± 0.006	0.483 ± 0.037
Ours (Paired)	0.676 ± 0.006	0.993 ± 0.012
<i>BTCV (external validation)</i>		
seg-only (LoRA)	0.745	0.312
Ours (Paired)	0.704	1.000

score and Switch Accuracy with McNemar test for statistical significance ($p < 0.001$).

Baselines. (1) **Frozen:** SAM-Med3D without adaptation; (2) **seg-only:** LoRA fine-tuning with standard Dice+BCE loss, single prompt per step; (3) **Ours:** Paired Prompt-Contrastive Training with full loss (Eq. 3).

4.2 Main Results

Table 1 presents the main comparison on AMOS (3 seeds) and BTCV (external validation).

Standard fine-tuning (seg-only) improves Dice substantially (0.450 $\rightarrow$ 0.720) but Switch Accuracy only reaches 0.483—the model still ignores prompts nearly half the time. Our method achieves Switch Accuracy 0.993 (+51.0 points over seg-only) with a 4.4-point Dice cost. The Dice gap is modest and expected: some segmentation capacity is redirected toward prompt-conditioned behavior rather than unconditional organ prediction.




Figure 3: Training dynamics over 50 epochs. (a) Switch Accuracy: our method maintains >0.95 throughout while seg-only stagnates at ~ 0.5 . (b) Dice: both methods improve, with seg-only reaching slightly higher final Dice due to unconstrained optimization.

On BTCV external validation, the improvement is even larger: Switch 1.000 vs 0.312 (+68.8 points), demonstrating strong generalization. The consistent pattern across datasets confirms that paired training fundamentally changes model behavior rather than overfitting to dataset-specific patterns.

4.3 Training Dynamics

Figure 3 shows Dice and Switch Accuracy over 50 training epochs. The seg-only baseline exhibits a characteristic pattern: Dice climbs steadily while Switch Accuracy fluctuates around 0.5 without improvement—the model learns better segmentation but not better prompt following. In contrast, our method maintains Switch Accuracy above 0.95 from epoch 1 onward while Dice gradually improves, demonstrating that paired training establishes prompt faithfulness immediately and maintains it throughout optimization.

4.4 Mechanism Ablation

To isolate the contribution of each component, we conduct a factor-isolation experiment (Table 2, 20 epochs, seed=42). We vary two binary factors—paired vs. single prompt training, and presence vs. absence of auxiliary losses—to determine what drives faithfulness. All configurations use identical training budgets: same number of optimizer steps, same images per epoch, and same LoRA parameters. Paired methods process two decoder passes per image (reusing the cached encoder embedding), while single-prompt methods process one. This means paired methods see twice as many prompt-target supervision signals per step; however, the seg-only baseline trained for equivalent total prompt exposures (100 epochs) still achieves only Switch 0.55, confirming that exposure alone is insufficient without simultaneous pairing.

The results reveal a clear hierarchy of importance:

Paired training is necessary and sufficient (+40 points). Moving from single-prompt to paired-prompt training—even without any auxiliary loss—jumps Switch Accuracy from 0.55 to 0.94. This is the dominant mechanism.

Grounding adds marginal benefit on top of paired training (+5 points, $0.94 \rightarrow 0.99$). However, grounding *without* paired training actually *decreases* Switch Accuracy ($0.55 \rightarrow 0.43$). This counterintuitive result suggests that grounding reinforces the model tendency to predict based

Table 2: Mechanism ablation (20 epochs, AMOS). Paired training is the dominant factor (+40pt Switch). Grounding without paired training is counterproductive.

Method	Paired	Sw. Acc	Dice
seg-only	No	0.55	0.62
+ grounding only	No	0.43	0.62
paired (seg only)	Yes	0.94	0.60
+ grounding	Yes	0.99	0.57
+ ground + switch	Yes	0.99	0.60




Figure 4: Mechanism ablation visualization. Single-prompt methods (red) fail regardless of loss design. Paired-prompt methods (green) succeed regardless of auxiliary losses. The gap between groups (~ 40 points) dwarfs within-group variation.

on spatial priors without learning to differentiate between prompts.

Switch margin loss adds negligible benefit ($0.99 \rightarrow 0.99$). When paired training already provides the contrastive signal, the explicit margin loss is redundant.

Figure 4 visualizes this hierarchy, making clear that the contribution is a *training paradigm* (paired exposure), not a specific loss function.

4.5 Hyperparameter Robustness

Table 3 demonstrates that all paired-training configurations maintain Switch Accuracy ≥ 0.978 across a wide range of margin values and loss weights.

The Dice score varies modestly with λ_{switch} (0.545–0.595), suggesting a mild tradeoff: higher switch weight slightly reduces segmentation accuracy. However, Switch Accuracy remains near-perfect across all settings, indicating that the paired training mechanism is inherently robust and does not require careful tuning.

4.6 Cross-Domain Analysis

To assess whether prompt unfaithfulness is specific to 3D medical SAM, we evaluate SAM2 [Ravi et al., 2024] on COCO natural images using the same Switch Accuracy protocol. SAM2 achieves 0.954 Switch Accuracy out of the box. While direct comparison is limited by differences in model scale, training data, and modality, this result suggests that prompt unfaithfulness is not inherent to the promptable segmentation paradigm but rather emerges under the specific conditions of 3D medical adaptation.

Table 3: Hyperparameter robustness (20 epochs, AMOS). All paired-training configurations maintain Switch ≥ 0.978 .

Configuration	Switch Acc.	Dice
margin = 0.1	0.990	0.561
margin = 0.2 (default)	0.990	0.595
margin = 0.3	0.990	0.555
$\lambda_{\text{switch}} = 0.25$	0.989	0.567
$\lambda_{\text{switch}} = 1.0$ (default)	0.990	0.595
$\lambda_{\text{switch}} = 2.0$	0.990	0.545




Figure 5: Qualitative comparison. Each row shows predictions for a different prompt target on the same image. seg-only produces identical outputs regardless of prompt (unfaithful). Our method correctly segments the prompted organ (faithful).

We hypothesize this specificity arises from two factors: (1) limited diversity in medical training data (anatomical structures are highly consistent across patients), enabling the model to exploit spatial priors; and (2) the frozen encoder paradigm, where only the decoder is adapted, limiting the model capacity to learn new prompt-conditioned features.

4.7 Qualitative Results

Figure 5 shows representative examples comparing seg-only and our method. When given two different organ prompts on the same image, seg-only produces nearly identical predictions (both targeting the dominant organ), while our method correctly switches the segmentation target to match each prompt.

5 Analysis and Discussion

5.1 Why Paired Training Works

The effectiveness of paired prompt training can be understood through the lens of counterfactual supervision. In single-prompt training, the model observes (x, p_A, y_A) and learns $f(x, p_A) \approx y_A$. However, nothing prevents the model from learning $f(x, \cdot) \approx y_A$ for *all* prompts—the training signal is consistent with prompt-independent behavior. Paired training adds the simultaneous constraint $f(x, p_B) \approx y_B$, which is only satisfiable if the model actually uses the prompt to differentiate its output.

This is analogous to the distinction between correlation and causation in causal inference: single-prompt training

establishes correlation between prompts and outputs (they co-occur), while paired training establishes a causal relationship (changing the prompt *causes* the output to change). The model must learn a prompt-conditioned mapping rather than a prompt-independent one.

Importantly, this mechanism is *architecture-agnostic*. It does not depend on specific model components or loss formulations—any promptable model trained with paired prompts will be forced to use the prompt signal. This generality is confirmed by our ablation: the specific auxiliary losses (grounding, switch margin) contribute minimally compared to the paired exposure itself.

5.2 Why Grounding Alone Fails

A surprising finding is that grounding loss *without* paired training actually decreases Switch Accuracy (0.55→0.43). We hypothesize this occurs because grounding reinforces the model’s spatial attention to the prompt location without requiring output differentiation. The model learns to activate the region around the prompt point, but since anatomical priors dominate, it still predicts the same organ—just with slightly better spatial localization of that prediction. In effect, grounding makes the model more confident in its unfaithful prediction rather than more faithful.

This finding has practical implications: naive attempts to improve prompt following by adding prompt-location supervision can backfire without the paired training paradigm. The model needs to see *contrasting* prompts, not just be told to attend to the prompt location.

5.3 The Dice-Switch Tradeoff

Our method incurs a 4.4-point Dice cost compared to seg-only. This tradeoff is inherent and can be understood quantitatively. A model that always predicts the largest organ achieves high average Dice (by being correct for the dominant class) but zero Switch Accuracy. Conversely, a perfectly faithful model must sometimes predict smaller, harder organs when prompted, reducing average Dice because small organs are inherently harder to segment.

The 4.4-point gap represents the cost of controllability—a modest price for transforming an uncontrollable model into a reliable interactive tool. In clinical practice, a model that achieves 0.72 Dice but ignores the clinician 52% of the time is far less useful than one achieving 0.68 Dice that reliably follows instructions 99% of the time.

We note that the absolute Dice of 0.676 is achieved on a challenging fixed prompt bank that includes all organ pairs (including small organs like adrenal glands). Per-organ analysis shows that large organs (liver, spleen) maintain Dice > 0.85 while the gap concentrates on small, anatomically challenging structures where both methods struggle.

5.4 Implications for Evaluation Practice

Our findings have broader implications for how interactive segmentation models should be evaluated. Standard practice reports only Dice (or IoU) averaged across prompts, which conflates segmentation quality with prompt faithfulness. We

Table 4: Computational overhead analysis. Paired training adds minimal cost during training and zero cost at inference.

Metric	seg-only	Ours
Total parameters	100.5M	100.5M
Trainable parameters	1.9M (1.89%)	1.9M (1.89%)
Training time/epoch	1117s	1177s (+5.3%)
Total training (50ep)	6.2h	6.5h
Inference time	identical	identical
Checkpoint size	7.6MB	7.6MB

argue that Switch Accuracy—or similar prompt-faithfulness diagnostics—should complement standard Dice evaluation for models claiming interactive or promptable segmentation capability.

The problem we identify is not unique to SAM-Med3D—it is a structural consequence of single-prompt training with strong image priors. Any promptable model trained on data where certain classes dominate (common in medical imaging) is susceptible to this failure mode. Our Switch Accuracy metric provides a simple, computationally cheap diagnostic that can be applied to any promptable segmentation model.

5.5 Computational Efficiency

Table 4 summarizes the computational overhead. The second decoder forward pass adds only 5.3% training time because the expensive image encoder computation (which dominates wall-clock time) is shared between both prompts. At inference, only a single prompt is processed per interaction (standard clinical use), so there is zero deployment overhead. The method requires no architectural changes, no additional parameters, and no modifications to the inference pipeline.

5.6 Limitations and Future Work

Single architecture. We validate on SAM-Med3D only. While the mechanism (paired exposure) is architecture-agnostic in principle, empirical validation on SegVol, Medical SAM2, or other 3D medical foundation models would strengthen generality claims. We attempted SegVol validation but encountered version incompatibility issues that prevented fair comparison.

Dice cost. The 4.4-point Dice gap, while modest, exists. Future work could explore Pareto-optimal training schedules (e.g., curriculum from paired to single-prompt) or adaptive loss weighting to minimize this cost while maintaining faithfulness.

Annotation requirement. Paired prompt sampling requires multi-organ annotations in the same volume. This is readily available in standard benchmarks (AMOS, BTCV, TotalSegmentator) but may limit applicability to single-organ datasets. Synthetic pairing strategies (e.g., using pseudo-labels for the second organ) could address this limitation.

Point prompts only. We evaluate point-prompt faithfulness. Extension to box prompts (MedSAM) and text prompts (SegVol) remains important future work, though the paired training principle should transfer directly.

Clinical validation. While our quantitative results are strong, clinical deployment would require user studies to assess whether improved Switch Accuracy translates to improved annotation efficiency and clinician satisfaction in practice.

6 Conclusion

We identify prompt unfaithfulness as a hidden failure mode in 3D medical Segment Anything Models: these models achieve acceptable segmentation Dice while largely ignoring user prompts (Switch Accuracy 20.8%). This failure is distinct from prompt robustness, hidden by standard evaluation metrics, and catastrophic for interactive clinical use.

We propose Switch Accuracy as a diagnostic metric that exposes this failure, and trace the root cause to single-prompt training, which never constrains counterfactual prompt behavior. Our solution—Paired Prompt-Contrastive Training—exposes the model to two different prompts per image during training, forcing prompt-conditioned output differentiation.

Comprehensive mechanism ablation reveals that paired prompt exposure is the dominant factor (+40 points Switch Accuracy), while specific loss designs provide only marginal additional benefit. The method achieves 99.3% Switch Accuracy (from 20.8% baseline) with 4.4-point Dice cost, generalizes across datasets (BTCV: Switch 1.000), and adds negligible overhead (5.3% training time, zero inference cost).

Our findings suggest that prompt faithfulness should be a standard evaluation criterion for interactive segmentation models, alongside traditional accuracy metrics. The paired training paradigm is simple, effective, and broadly applicable to any promptable model where controllability is essential.

References

- Soumitri Chattopadhyay, Basar Demir, and Marc Niethammer. On the robustness of foundational 3d medical image segmentation models against imprecise visual prompts. *arXiv preprint arXiv:2601.16383*, 2026.
- Cheng Chen, Juzheng Miao, Dufan Wu, Zhiling Yan, Sekeun Kim, Jiang Hu, Aoxiao Zhong, Zhengliang Liu, Lichao Sun, Xiang Li, and Quanzheng Li. MA-SAM: Modality-agnostic SAM adaptation for 3d medical image segmentation. *Medical Image Analysis*, 98:103310, 2024a.
- Tianrun Chen, Ankang Lu, Lanyun Zhu, Chaotao Ding, Chunan Yu, Deyi Ji, Zejian Li, Lingyun Sun, Papa Mao, and Ying Zang. SAM2-Adapter: Evaluating and adapting segment anything 2 in downstream tasks. *arXiv preprint arXiv:2408.04579*, 2024b.
- Wangang Cheng, Junfei Li, and Yuting Wang. Tensor CUR decomposition based low-rank parameter adaptation and its application in medical image segmentation. *arXiv preprint arXiv:2501.02227*, 2025.
- Xingyu Deng, Lei Zhu, and Hao Chen. Few-shot adaptation of training-free foundation model for 3d medical image segmentation. *arXiv preprint arXiv:2501.09138*, 2025.

- Yuxin Du, Fan Bai, Tiejun Huang, and Bo Zhao. SegVol: Universal and interactive volumetric medical image segmentation. *arXiv preprint arXiv:2311.13385*, 2023.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shanen Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *ICLR*, 2022.
- Yao Huang, Xingru Li, Zhongyu Chen, and Yunhe Gao. Perturbed prompts for robust adaptation of segment anything model for polyp segmentation. *arXiv preprint arXiv:2405.16740*, 2024.
- Yuanfeng Ji, Haotian Bai, Chongjian Ge, Jie Yang, Ye Zhu, Ruimao Zhang, Zhen Li, Lingyan Zhang, Wanling Ma, Xiang Wan, et al. AMOS: A large-scale abdominal multi-organ benchmark for versatile medical image segmentation. In *NeurIPS Datasets and Benchmarks*, 2022.
- Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. In *ICCV*, 2023.
- Bennett Landman, Zhoubing Xu, Juan Eugenio Iglesias, Martin Styner, Thomas Langerak, and Arno Klein. Multi-atlas labeling beyond the cranial vault. *MICCAI Workshop*, 2015.
- Mengxi Li, Hao Chen, Yue Zhang, and Liansheng Wang. Saliency-guided prompt distillation for robust segmentation with SAM. *arXiv preprint arXiv:2604.23314*, 2024.
- Xinrui Liu, Spyridon Thermos, Alison O’Neil, and Sotirios A Tsaftaris. Overcoming dimensional collapse in self-supervised contrastive learning for medical image segmentation. *arXiv preprint arXiv:2402.14611*, 2024.
- Jun Ma, Yuting He, Feifei Li, Lin Han, Chenyu You, and Bo Wang. Segment anything in medical images. *Nature Communications*, 15:654, 2024.
- Nikhila Ravi, Valentin Gabeur, Yuan-Ting Hu, Ronghang Hu, Chaitanya Ryali, Tengyu Ma, Haitham Khedr, Roman Rädle, Chloe Rolland, Laura Gustafson, et al. SAM 2: Segment anything in images and videos. *arXiv preprint arXiv:2408.00714*, 2024.
- Haoyu Wang, Sizheng Guo, Jin Ye, Zhongying Deng, Junlong Cheng, Tianbin Li, Jianpin Chen, Yanzhou Su, Ziyang Huang, Yiqing Shen, et al. SAM-Med3D. *arXiv preprint arXiv:2310.15161*, 2023.
- Yizhou Wang, Liang Zhang, and Hao Chen. Domain and task-focused example selection for data-efficient contrastive medical image segmentation. *arXiv preprint arXiv:2505.19208*, 2025.
- Zhihao Wu, Yuliang Shi, Shengjie Huang, and Zhongxing Ding. LoRA-MedSAM: Efficient medical image segmentation. In *Springer LNCS*, 2024.
- Yuxuan Zhao, Xiang Li, and Hao Chen. A survey of contrastive learning in medical AI: Foundations, biomedical modalities, and future directions. *Preprints*, 2024.
- Jiayuan Zhu, Abdullah Hamdi, Yunli Qi, Yueming Jin, and Junde Wu. Medical SAM 2: Segment medical images as video via segment anything model 2. *arXiv preprint arXiv:2408.00874*, 2024.